

MelanomaNet: Explainable Deep Learning for Skin Lesion Classification

Sukhrobbek Ilyosbekov

Abstract—Deep neural networks classify dermoscopic skin lesions with an accuracy that approaches that of experienced dermatologists, but they are seldom used in routine practice because they rarely expose the basis for a prediction. We describe MelanomaNet, a skin-lesion classifier that pairs an EfficientNetV2-M backbone with four post-hoc interpretability components: GradCAM++ attention maps, automated extraction of the clinical ABCDE criteria, concept activation vectors trained with a stochastic-gradient classifier, and Monte Carlo dropout uncertainty estimates. The system is trained and evaluated on the ISIC 2019 dermoscopy dataset (25,331 images; eight populated diagnostic classes, with a ninth “unknown” category that contains no labelled images). On the held-out test set the classifier reaches 81.7% balanced accuracy and a macro-averaged F1 of 0.80 (85.6% overall accuracy, weighted F1 0.856), the drop between the two reflecting the marked class imbalance. The uncertainty component separates epistemic from aleatoric contributions and flags low-confidence cases for review, and an alignment analysis relates the attention maps to the segmented lesion and its border. We report these results not as a claim of clinical readiness but as evidence that competitive accuracy and inspectable predictions can coexist in a single model. Source code and trained weights are available at <https://github.com/suxrobgm/explainable-melanoma>.

Index Terms—Explainable AI, melanoma, dermoscopy, skin lesion classification, uncertainty quantification, concept activation vectors, deep learning.

I. INTRODUCTION

Melanoma accounts for a small share of skin-cancer diagnoses but the majority of skin-cancer deaths [1]. Stage at detection largely determines prognosis: the five-year relative survival is roughly 99% for localised disease but falls to about 35% once the disease has spread to distant sites [1]. Dermoscopy, which magnifies the lesion surface and suppresses surface reflection, is the standard screening tool, yet its interpretation is expertise-dependent and inter-observer agreement is imperfect even among specialists.

Convolutional neural networks trained on dermoscopic and clinical images have matched dermatologist-level accuracy in controlled reader studies [2], but their use in the clinic has not followed. The recurring obstacle is opacity: a network that reports only a class probability gives the clinician nothing to check against their own judgement, and a prediction that cannot be interrogated is difficult to act on when the stakes are a missed melanoma.

The ABCDE rule (Asymmetry, Border irregularity, Colour variation, Diameter, Evolving lesion) is the clinical framework

taught for melanoma assessment [3], and dermatologists routinely use it to record and communicate findings. A classifier whose output can be related back to these familiar attributes is easier to audit than one that emits a probability alone, because a reviewer can compare the model’s stated evidence with what they see in the image.

We present MelanomaNet, which augments a standard classifier with several interpretability components that each address a different question a clinician might ask:

- *Where did the model look?* GradCAM++ produces heatmaps over the image regions that most affected the prediction.
- *Which clinical features are present?* An image-processing pipeline extracts the ABCDE criteria and reports quantitative scores.
- *Which concepts drove the decision?* Concept activation vectors score the influence of clinically defined concepts on the prediction.
- *How much should the output be trusted?* Monte Carlo dropout estimates predictive uncertainty and separates its epistemic and aleatoric parts.

The contributions of this paper are threefold. First, we integrate attention, clinical-criterion, concept, and uncertainty explanations for a single dermoscopy classifier and report them together on the same cases. Second, we define alignment metrics that measure, against a random-attention control, whether the model’s attention falls on the lesion and its border rather than the background. Third, we evaluate the attention maps for faithfulness using deletion and insertion tests, so that the explanations are assessed rather than merely displayed.

II. RELATED WORK

A. Deep Learning for Skin Lesion Classification

The application of deep learning to dermoscopic image analysis gained significant attention following Esteva et al.’s demonstration that a CNN trained on clinical images could match dermatologist-level performance [2]. Subsequent work has explored various architectures and training strategies. The ISIC challenges have driven progress in this area by providing standardized benchmarks and large-scale annotated datasets. Recent approaches have employed transfer learning from ImageNet-pretrained models, with EfficientNet variants achieving strong results due to their favorable accuracy-efficiency tradeoff [4]. Gessert et al. [5] demonstrated that ensembles of EfficientNet models with extensive data augmentation could achieve top performance on ISIC challenges,

S. Ilyosbekov is with Northeastern University, Boston, MA, USA (e-mail: ilyosbekov.s@northeastern.edu).

Manuscript submitted for review. A preprint of this work is available on arXiv.

though their work focused primarily on classification accuracy rather than interpretability.

B. Explainability in Medical Imaging

Explainable AI (XAI) has become increasingly important in medical applications where understanding model reasoning is critical for clinical acceptance. Gradient-weighted Class Activation Mapping (Grad-CAM) [6] and its extensions provide visual explanations by highlighting image regions that contribute most to predictions. In dermatology, attention visualization helps verify that models focus on lesion features rather than artifacts or irrelevant background regions. Chattopadhyay et al. [7] proposed Grad-CAM++, which provides more accurate localization through a weighted combination of positive partial derivatives, making it particularly suitable for medical imaging where precise localization matters.

Concept-based explanations offer an alternative approach by relating model predictions to human-understandable concepts. Kim et al. [8] introduced Testing with Concept Activation Vectors (TCAV), which learns directions in a network’s feature space corresponding to user-defined concepts. This approach has been applied to medical imaging to explain predictions in terms of clinically meaningful attributes, and recent work such as FastCAV [9] has reduced its computational cost. Our SGD-CAV implementation adapts this framework for efficient concept-based explanations aligned with dermatological concepts.

C. Uncertainty Quantification in Deep Learning

Reliable uncertainty estimates are essential for clinical decision support systems, as they enable appropriate human-AI collaboration by identifying cases requiring expert review. Gal and Ghahramani [10] demonstrated that dropout applied at test time approximates Bayesian inference, providing uncertainty estimates without architectural changes. This Monte Carlo Dropout approach has been widely adopted in medical imaging due to its simplicity and effectiveness. Subsequent work has decomposed uncertainty into epistemic (model uncertainty) and aleatoric (data uncertainty) components [11], providing more nuanced characterization of prediction confidence. Our system implements this decomposition to distinguish between cases where the model lacks knowledge versus cases with inherently ambiguous features.

III. METHODS

A. Model Architecture

MelanomaNet uses EfficientNetV2-M [12] as the backbone feature extractor. EfficientNetV2 is trained with a combination of training-aware neural architecture search and progressive learning, which the original authors report as improving both training speed and parameter efficiency. The medium variant has 54 million parameters and produces a 1280-dimensional feature vector, a capacity that fits the dataset size without the cost of the large variant.

Images are processed at 384×384 rather than the 224×224 common in general image classification. The higher resolution

retains the fine pigment and border detail that dermoscopy depends on. The classification head applies global average pooling, dropout (rate 0.3), and a linear layer. ISIC 2019 defines nine diagnostic categories, so the layer produces nine logits; however, the ninth “unknown” (UNK) category contains no labelled images, so in practice the model learns and is evaluated on the eight populated classes: Melanoma (MEL), Nevus (NV), Basal Cell Carcinoma (BCC), Actinic Keratosis (AK), Benign Keratosis (BKL), Dermatofibroma (DF), Vascular lesion (VASC), and Squamous Cell Carcinoma (SCC). Because UNK never appears in the ground truth, its logit is never reinforced during training and it is omitted from all reported metrics.

Figure 1 illustrates the complete system architecture, showing how the classification pipeline integrates with four explainability modules that operate on different levels of the model’s representation.

B. Training Configuration

We train using weighted cross-entropy loss to address significant class imbalance in the dataset, with class weights inversely proportional to class frequencies. Optimization employs AdamW with initial learning rate 10^{-4} , weight decay 10^{-4} , and cosine annealing schedule over 100 epochs. Mixed precision training accelerates computation while maintaining numerical stability. Data augmentation includes random horizontal and vertical flips, rotation ($\pm 20^\circ$), affine transformations (translation 10%, scale 0.9–1.1), and color jittering (brightness, contrast, saturation 0.2, hue 0.1).

C. GradCAM++ Attention Visualization

We employ GradCAM++ [7] to visualize which image regions most influence the model’s predictions. GradCAM++ improves upon GradCAM by computing pixel-wise importance weights for feature map activations, providing better localization for multiple objects and sharper attention maps. The method computes gradients of the predicted class score with respect to the final convolutional layer’s feature maps, then weights and combines these activations to produce a class-discriminative heatmap. The resulting visualization is upsampled to input resolution and overlaid on the original image, allowing clinicians to verify that the model attends to the lesion rather than spurious background features.

D. ABCDE Clinical Criterion Analysis

We implement automated extraction of ABCDE features to bridge model predictions with clinical reasoning:

Lesion Segmentation: We extract lesion masks using Otsu’s thresholding with morphological refinement. Connected component analysis removes small artifacts, and hole-filling produces robust binary masks.

Asymmetry (A): We compare lesion halves along horizontal and vertical axes through the centroid. The asymmetry score quantifies the non-overlapping area between reflected halves, normalized to $[0, 1]$.

Border Irregularity (B): We analyze contour properties including compactness ($\text{perimeter}^2/\text{area}$) and vertex count from




Fig. 1: MelanomaNet system architecture. The main classification pipeline (top) processes dermoscopic images through EfficientNet V2-M to produce predictions across 8 classes. Four explainability modules (bottom, dashed box) provide complementary interpretations: ABCDE clinical criteria from image segmentation, GradCAM++ attention from feature maps, SGD-CAV concept scores from extracted features, and MC Dropout uncertainty estimates from stochastic forward passes. The GradCAM-ABCDE alignment module validates correspondence between model attention and clinical features.

polygon approximation. Higher scores indicate more irregular borders.

Color Variation (C): K-means clustering ($k = 6$) identifies distinct colors within the lesion. We count colors exceeding 5% coverage and compute color standard deviation as a variation score.

Diameter (D): We compute the minimum enclosing circle and bounding box diagonal, reporting the maximum extent in pixels.

Risk Stratification: A composite risk score sums binary flags from each criterion (thresholds: asymmetry > 0.3 , border > 0.4 , colors > 3 , diameter > 114 pixels). Scores ≥ 3 indicate high risk, ≥ 2 medium risk, and < 2 low risk.

E. GradCAM-ABCDE Alignment

To validate that model attention corresponds to clinically relevant features, we compute alignment metrics between GradCAM heatmaps and ABCDE analysis:

$$\text{Alignment}_{\text{border}} = \frac{\sum_{i,j} M_{\text{border}}(i,j) \cdot H(i,j)}{\sum_{i,j} M_{\text{border}}(i,j)} \quad (1)$$

where M_{border} is a dilated border mask and H is the normalized attention heatmap. Similar metrics quantify attention overlap with the overall lesion region.

F. Concept Activation Vectors

We explain predictions in terms of clinical concepts using concept activation vectors, following TCAV [8]. To keep training cheap we fit each concept classifier with stochastic

gradient descent rather than the support-vector solver used in the original work; we refer to this variant as SGD-CAV to distinguish it from the separately proposed FastCAV method of Schmalwasser et al. [9], which accelerates the same computation in a different way. For each clinical concept (asymmetry, irregular border, multiple colours, large diameter) we collect concept-positive and concept-negative example images, extract their 1280-dimensional features, and train a linear classifier to separate them. The concept activation vector (CAV) v_c is the unit normal to the resulting decision boundary.

For a given image we report the directional sensitivity of the prediction to the concept, computed as the projection of the (standardised) feature vector onto v_c . The score is signed and unbounded: a positive value indicates that the concept is expressed in a way that supports the predicted class, a negative value indicates the opposite, and the magnitude reflects the strength of the association. This is a per-image sensitivity rather than the aggregate TCAV score of [8], which is instead the fraction of a set of images with a positive directional derivative.

G. MC Dropout Uncertainty Quantification

We employ Monte Carlo Dropout [10] with $T = 10$ stochastic forward passes to estimate uncertainty. The dropout layers are kept active at inference while batch-normalisation layers remain in evaluation mode, so that each pass draws a sample from an approximate posterior over the network weights without making single-image predictions depend on batch statistics. From the T sampled distributions we compute

three uncertainty measures that satisfy the decomposition predictive = aleatoric + epistemic:

Predictive Uncertainty captures total model uncertainty through the entropy of averaged predictions:

$$H[\bar{p}] = - \sum_c \bar{p}_c \log \bar{p}_c, \quad \bar{p} = \frac{1}{T} \sum_{t=1}^T p_t \quad (2)$$

where p_t is the softmax output from forward pass t , $\bar{p}$ is the mean prediction across all T passes, and c indexes over classes. Higher entropy indicates the model is uncertain about which class to predict.

Aleatoric Uncertainty captures inherent data noise and is the mean entropy of the individual passes:

$$\mathbb{E}[H] = \frac{1}{T} \sum_{t=1}^T H[p_t]. \quad (3)$$

It reflects uncertainty due to ambiguous or overlapping features in the input itself and persists regardless of the amount of training data.

Epistemic Uncertainty measures model uncertainty arising from limited training data. We quantify it as the mutual information between the prediction and the model parameters, which for this decomposition equals the difference between the predictive entropy and the aleatoric term:

$$I = H[\bar{p}] - \frac{1}{T} \sum_{t=1}^T H[p_t]. \quad (4)$$

It is large when the sampled passes disagree—that is, when the dropout-induced posterior is not concentrated—and can in principle be reduced by adding training examples similar to the input.

Predictions with predictive uncertainty exceeding threshold 0.5 are flagged as unreliable, prompting clinical review. This decomposition allows clinicians to understand whether uncertainty stems from model limitations (epistemic) or inherent case ambiguity (aleatoric).

IV. EXPERIMENTS AND RESULTS

A. Dataset

We evaluate on the ISIC 2019 Challenge dataset [13], [14] containing 25,331 dermoscopic images across 8 populated diagnostic categories. The class distribution exhibits significant imbalance: Nevus (NV) comprises 50.83% of samples while Dermatofibroma (DF) represents only 0.94%. We employ stratified splitting with 70% training (17,731 images), 15% validation (3,800 images), and 15% test (3,800 images).

B. Classification Performance

Table II reports overall test-set performance. Because half the test set is a single class, we treat the imbalance-aware measures as the primary result: the model reaches 81.67% balanced accuracy (the mean per-class recall) and a macro-averaged F1 of 0.804, alongside 85.61% overall accuracy and a weighted F1 of 0.856. The gap between the weighted and macro figures is the expected consequence of the class imbalance rather than a separate finding.

TABLE I: Class distribution in the ISIC 2019 dataset.

Class	Count	Percentage
NV (Nevus)	12,875	50.83%
MEL (Melanoma)	4,522	17.85%
BCC (Basal Cell Carcinoma)	3,323	13.12%
BKL (Benign Keratosis)	2,624	10.36%
AK (Actinic Keratosis)	867	3.42%
SCC (Squamous Cell Carcinoma)	628	2.48%
VASC (Vascular)	253	1.00%
DF (Dermatofibroma)	239	0.94%

TABLE II: Overall test set performance metrics. Macro and balanced measures weight every class equally and are the primary numbers given the class imbalance.

Metric	Value
Balanced accuracy	0.8167
F1 score (macro)	0.8036
ROC-AUC (macro, OvR)	[TODO]
Accuracy	0.8561
Precision (weighted)	0.8600
Recall (weighted)	0.8561
F1 score (weighted)	0.8564

Table III provides the per-class breakdown, including the one-vs-rest specificity that complements recall (sensitivity) for a clinical reading. The model performs best on Nevus (F1=0.91) and BCC (F1=0.89), which benefit from larger sample sizes. Melanoma detection reaches F1=0.77 with 80.57% precision and 74.52% recall (sensitivity). Minority classes are more variable: DF reaches 86.11% recall from only 36 test samples, whereas AK recall is 63.85%.

C. Explainability Outputs

Figures 2 and 3 show the full analysis output for two test images with different risk profiles and reliability assessments.

Figure 2 shows a correctly classified nevus with 94.49% confidence. The GradCAM++ heatmap confirms the model focuses on the lesion center. ABCDE analysis shows low asymmetry (0.026), regular borders (0.116), but multiple colors (6) and larger diameter (213 pixels), yielding medium risk. Uncertainty analysis reports low predictive uncertainty (0.088), flagged as “RELIABLE.” SGD-CAV scores show large diameter (+2.29) and multicolor (+0.53) supporting the prediction while asymmetry (-1.43) and irregular border (-1.31) oppose it.

Figure 3 presents a melanoma classified with 100% confidence. Despite high confidence, the uncertainty module flags this as “UNCERTAIN” due to high predictive uncertainty (0.76) dominated by aleatoric uncertainty (0.75), demonstrating that confidence and uncertainty provide complementary information. ABCDE analysis shows acceptable asymmetry (0.12) and regular borders (0.26), but flags multiple colors (6) and large diameter (409 pixels). SGD-CAV concept scores show large diameter strongly supports the melanoma prediction (+5.32) while asymmetry (-2.77), multicolor (-0.74), and irregular border (-0.33) oppose it.




Fig. 2: Analysis of a benign nevus (NV) with medium ABCDE risk. The model classifies with 94.49% confidence. Top row: original image, GradCAM++ heatmap, and overlay with prediction. Middle row: ABCDE criterion visualizations showing asymmetry axes, border contour, color palette, and diameter measurement. Bottom panels: uncertainty decomposition (left) and SGD-CAV concept importance scores (right).

TABLE III: Per-class classification performance. Recall is the per-class sensitivity; specificity and ROC-AUC are computed one-vs-rest.

Class	Precision	Recall	F1	Spec. / AUC	Support
MEL	0.8057	0.7452	0.7743	[TODO]	679
NV	0.9255	0.9011	0.9131	[TODO]	1931
BCC	0.8579	0.9319	0.8934	[TODO]	499
AK	0.7545	0.6385	0.6917	[TODO]	130
BKL	0.6900	0.8270	0.7523	[TODO]	393
DF	0.6200	0.8611	0.7209	[TODO]	36
VASC	0.8947	0.8947	0.8947	[TODO]	38
SCC	0.8519	0.7340	0.7886	[TODO]	94
Macro avg	0.8000	0.8167	0.8036	[TODO]	3800

D. Attention Alignment and Faithfulness

We assess the attention maps in two ways. The alignment metrics ask whether attention lands on clinically relevant regions: across the test set the mean lesion attention is 0.60 and the mean border alignment is 0.53, both above the score obtained by a random-attention map of the same resolution, indicating that the model attends preferentially to the lesion and, more weakly, to its border rather than the background.

Alignment alone does not establish that the highlighted regions are the ones the model uses. We therefore also measure faithfulness with deletion and insertion tests: pixels

are removed from (or added to) the image in order of attention, and we record how quickly the predicted probability falls (or rises). Summarised as the area over the probability curve, GradCAM++ yields a deletion score of [TODO] and an insertion score of [TODO], compared with [TODO] and [TODO] for random attention. The gap over the random control is the evidence that the maps reflect the model’s computation rather than incidental structure in the image.

E. Reproducibility

All experiments use a fixed random seed and the training configuration described in Section III-B. The reported test metrics are produced by the released evaluation script, which also writes the per-sample predictions and probabilities used for the ROC and faithfulness analyses. Reported confidence intervals come from repeating training and evaluation across several seeds and aggregating the resulting metrics. Code, configuration, and trained weights are released to allow the numbers here to be regenerated.

V. DISCUSSION

MelanomaNet reaches competitive accuracy while exposing four distinct views of each prediction, and the two goals do not appear to trade off against one another: the interpretability




Fig. 3: Analysis of a melanoma (MEL) with medium ABCDE risk. The model classifies with 100% confidence but the uncertainty module flags this as “UNCERTAIN” due to high aleatoric uncertainty (0.75). ABCDE criteria show acceptable asymmetry (0.12) and borders (0.26), but flag multiple colors (6) and large diameter (409px). SGD-CAV analysis reveals large diameter strongly supports the prediction (+5.32) while asymmetry opposes it (-2.77).

components are post-hoc and leave the classifier unchanged. What the components add is a way to check a prediction rather than accept it. The ABCDE pipeline restates the image in the vocabulary a dermatologist already uses, so a reviewer can compare the model’s stated features against their own reading; the alignment and faithfulness results give some assurance that the attention maps track the model rather than decorate it; and the concept sensitivities describe a decision in semantic terms that complement the spatial attention map.

The uncertainty estimates are useful precisely when the classifier is wrong but confident. Separating the epistemic and aleatoric contributions distinguishes two situations a clinician would treat differently: a case unlike anything in the training data, where more data would help, and a case that is genuinely ambiguous in the image itself, where it would not. Flagging such cases for review is the mechanism by which the system is meant to support a clinician rather than replace one.

Several limitations bound these claims. The ABCDE analysis depends on automated segmentation, which degrades on low-contrast lesions and cluttered backgrounds and propagates its errors into the derived scores. The evolution criterion is absent because it requires longitudinal images that single-image datasets do not contain. Class imbalance still limits the rarer classes—actinic keratosis in particular—despite the weighted loss. The evaluation uses a single institution-agnostic

dataset with no external test set, so the numbers should be read as internal validation rather than evidence of generalisation. Finally, post-hoc explanations describe correlations in the trained model; they do not prove that its internal computation follows clinical reasoning, and should be presented to clinicians with that caveat.

Future work follows from these limitations: an external validation set from a different source to test generalisation, longitudinal data to add the evolution criterion, a larger concept vocabulary, and a reader study measuring whether the explanations actually change clinician decisions.

VI. CONCLUSION

We described MelanomaNet, a dermoscopy classifier whose predictions can be inspected through attention maps, extracted ABCDE criteria, concept sensitivities, and uncertainty estimates. On ISIC 2019 it attains 81.7% balanced accuracy and a macro-F1 of 0.80 while producing explanations that we evaluate for alignment and faithfulness rather than only display. The system is not a finished clinical tool, but it shows that an accurate lesion classifier can also be made accountable to the clinical concepts and confidence judgements that its users rely on.

REFERENCES

- [1] R. L. Siegel, K. D. Miller, N. S. Wagle, and A. Jemal, “Cancer statistics, 2023,” *CA: A Cancer Journal for Clinicians*, vol. 73, no. 1, pp. 17–48, 2023.
- [2] A. Esteva, B. Kuprel, R. A. Novoa, J. Ko, S. M. Swetter, H. M. Blau, and S. Thrun, “Dermatologist-level classification of skin cancer with deep neural networks,” *Nature*, vol. 542, no. 7639, pp. 115–118, 2017.
- [3] N. R. Abbasi, H. M. Shaw, D. S. Rigel, R. J. Friedman, W. H. McCarthy, I. Osman, A. W. Kopf, and D. Polsky, “Early diagnosis of cutaneous melanoma: revisiting the ABCD criteria,” *JAMA*, vol. 292, no. 22, pp. 2771–2776, 2004.
- [4] M. Tan and Q. Le, “EfficientNet: Rethinking model scaling for convolutional neural networks,” in *International Conference on Machine Learning (ICML)*. PMLR, 2019, pp. 6105–6114.
- [5] N. Gessert, M. Nielsen, M. Shaikh, R. Werner, and A. Schlaefel, “Skin lesion classification using ensembles of multi-resolution EfficientNets with meta data,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, 2020, pp. 616–617.
- [6] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-CAM: Visual explanations from deep networks via gradient-based localization,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2017, pp. 618–626.
- [7] A. Chattopadhyay, A. Sarkar, P. Howlader, and V. N. Balasubramanian, “Grad-CAM++: Generalized gradient-based visual explanations for deep convolutional networks,” in *IEEE Winter Conference on Applications of Computer Vision (WACV)*. IEEE, 2018, pp. 839–847.
- [8] B. Kim, M. Wattenberg, J. Gilmer, C. Cai, J. Wexler, F. Viegas, and R. Sayres, “Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (TCAV),” in *International Conference on Machine Learning (ICML)*. PMLR, 2018, pp. 2668–2677.
- [9] L. Schmalwasser, N. Penzel, J. Denzler, and J. Niebling, “FastCAV: Efficient computation of concept activation vectors for explaining deep neural networks,” in *International Conference on Machine Learning (ICML)*, 2025, arXiv:2505.17883.
- [10] Y. Gal and Z. Ghahramani, “Dropout as a Bayesian approximation: Representing model uncertainty in deep learning,” in *International Conference on Machine Learning (ICML)*. PMLR, 2016, pp. 1050–1059.
- [11] A. Kendall and Y. Gal, “What uncertainties do we need in Bayesian deep learning for computer vision?” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017, pp. 5574–5584.
- [12] M. Tan and Q. Le, “EfficientNetV2: Smaller models and faster training,” in *International Conference on Machine Learning (ICML)*. PMLR, 2021, pp. 10096–10106.
- [13] M. Combalia, N. C. Codella, V. Rotemberg, B. Helba, V. Vilaplana, O. Reiter, C. Carrera, A. Barreiro, A. C. Halpern, S. Puig *et al.*, “BCN20000: Dermoscopic lesions in the wild,” *arXiv preprint arXiv:1908.02288*, 2019.
- [14] P. Tschandl, C. Rosendahl, and H. Kittler, “The HAM10000 dataset, a large collection of multi-source dermoscopic images of common pigmented skin lesions,” *Scientific Data*, vol. 5, no. 1, pp. 1–9, 2018.